

Assignment #7 BigData – Spark | Name: Akshay Jain | CWID: A20502846

Exercise 1) Step A

Start up a Hadoop cluster as previously, but instead of choosing the “Core Hadoop” configuration chose the “Spark” configuration (see below), otherwise proceed as before.

General Configuration

Cluster name

☒ Logging ⓘ

S3 folder

Launch mode ☒ Cluster ⓘ ☐ Step execution ⓘ

Software configuration

Release ⓘ

Applications

- ☐ Core Hadoop: Hadoop 2.10.1, Hive 2.3.9, Hue 4.10.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.9.2
- ☐ HBase: HBase 1.4.13, Hadoop 2.10.1, Hive 2.3.9, Hue 4.10.0, Phoenix 4.14.3, and ZooKeeper 3.4.14
- ☐ Presto: Presto 0.267 with Hadoop 2.10.1 HDFS and Hive 2.3.9 Metastore
- ☒ Spark: Spark 2.4.8 on Hadoop 2.10.1 YARN and Zeppelin 0.10.0

Step B

```
EEEEEEEEEEEEEEEEEEEE MMMMMMM MMMMMMM RRRRRRRRRRRRRR
E:EEEEEEEEEEEEEEEE M:EEEE:M M:EEEE:M R:EEEEEEEE:R
EE:EEEEEEEEEEEE:E M:EEEE:M M:EEEE:M R:RRRRRR:R
E:EE EEEEE M:EEEE:M M:EEEE:M RR:R R:R
E:EE M:EE:M:M M:EE:M R:R R:R
E:EEEEEEEEEEEE M:EE:M M:EE:M M:EE:M R:RRRRRR:R
E:EEEEEEEEEEEE M:EE:M M:EE:M M:EE:M R:RRRRRR:R
E:EE M:EE:M M:EE:M M:EE:M R:R R:R
E:EE EEEEE M:EE:M MMM M:EE:M R:R R:R
EE:EEEEEEEE:E M:EE:M M:EE:M R:R R:R
E:EEEEEEEEEEEE:M M:EE:M RR:R R:R
EEEEEEEEEEEEEEEE MMMMMMM MMMMMMM RRRRRRR RRRRRR

[hadoop@ip-172-31-31-84 ~]$ hadoop fs -ls /user
Found 5 items
drwxrwxrwx - hadoop hdfsadmingroup 0 2022-10-23 03:39 /user/hadoop
drwxrwxrwx - livy livy 0 2022-10-23 03:39 /user/livy
drwxrwxrwx - root hdfsadmingroup 0 2022-10-23 03:39 /user/root
drwxrwxrwx - spark spark 0 2022-10-23 03:39 /user/spark
drwxrwxrwx - zeppelin hdfsadmingroup 0 2022-10-23 03:39 /user/zeppelin
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -mkdir /user/csp554
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -ls /user/
Found 6 items
drwxr-xr-x - hadoop hdfsadmingroup 0 2022-10-23 03:47 /user/csp554
drwxrwxrwx - hadoop hdfsadmingroup 0 2022-10-23 03:39 /user/hadoop
drwxrwxrwx - livy livy 0 2022-10-23 03:39 /user/livy
drwxrwxrwx - root hdfsadmingroup 0 2022-10-23 03:39 /user/root
drwxrwxrwx - spark spark 0 2022-10-23 03:39 /user/spark
drwxrwxrwx - zeppelin hdfsadmingroup 0 2022-10-23 03:39 /user/zeppelin
```

Use the TestDataGen program from previous assignments to generate new data files.
Copy both generated files to the HDFS directory “/user/hadoop”

```
[hadoop@ip-172-31-31-84 ~]$ java TestDataGen
Magic Number = 213024
```

Magic Number =213024

```
[hadoop@ip-172-31-31-84 ~]$ java TestDataGen
Magic Number = 213024
[hadoop@ip-172-31-31-84 ~]$ ls
foodplaces213024.txt  foodratings213024.txt  hql.zip  TestDataGen.class
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -cp foodratings213024.txt /user/csp554
cp: `foodratings213024.txt': No such file or directory
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -put foodratings213024.txt /user/csp554
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -put foodplaces213024.txt /user/csp554
[hadoop@ip-172-31-31-84 ~]$ ls
foodplaces213024.txt  foodratings213024.txt  hql.zip  TestDataGen.class
[hadoop@ip-172-31-31-84 ~]$ hadoop fs -ls /user/csp554
Found 2 items
-rw-r--r-- 1 hadoop hdfsadmin group 59 2022-10-23 03:54 /user/csp554/foodplaces213024.txt
-rw-r--r-- 1 hadoop hdfsadmin group 17447 2022-10-23 03:54 /user/csp554/foodratings213024.txt
[hadoop@ip-172-31-31-84 ~]$ pyspark
Python 3.7.10 (default, Jun 3 2021, 00:02:01)
[GCC 7.3.1 20180712 (Red Hat 7.3.1-13)] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
22/10/23 03:56:13 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.
Welcome to

  ____      _
 / ___|  __| | | |
| |  | |__| | | |
| |  | |__| | | |
| |  | |__| | | |
|_|  |____|_|_|_|

 version 2.4.8-amzn-2

Using Python version 3.7.10 (default, Jun 3 2021 00:02:01)
SparkSession available as 'spark'.
```

Step C

Load the 'foodratings' file as a 'csv' file into a DataFrame called foodratings. When doing so specify a schema having fields of the following names and types:

```
hadoop fs -copyFromLocal /home/hadoop/foodratings213024.txt
```

```
from pyspark.sql.types import *
struct1 = StructType().add("name", StringType(), True).add("food1", IntegerType(),
True).add("food2", IntegerType(), True).add("food3", IntegerType(), True).add("food4", IntegerType(),
True).add("placeid", IntegerType(), True) foodratings =
spark.read.schema(struct1).csv('foodratings213024.txt')
foodratings.printSchema()
foodratings.show(5)
```

```
>>> from pyspark.sql.types import *
>>> tab1=StructType().add("name",StringType(),True).add("food1",IntegerType(),True).add("food2",IntegerType(),True).add("food3",IntegerType(),True).add("food4",IntegerType(),True).add("placeid",IntegerType(),True)
>>> foodratings=spark.read.schema(tab1).csv('hdfs:///user/csp554/foodratings213024.txt')
```

```
>>> foodratings.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings.show(5)
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
| Joe|   13|    8|   15|   46|     5|
| Mel|    2|   49|   31|   47|     4|
| Joe|   23|   22|   22|    7|     5|
| Joe|   48|   26|   18|   35|     2|
| Mel|    9|   21|   40|   45|     3|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Exercise 2

Load the 'foodplaces' file as a 'csv' file into a DataFrame called foodplaces. When doing so specify a schema having fields of the following names and types

```
hadoop fs -copyFromLocal /home/hadoop/foodplaces213024.txt
```

```
from pyspark.sql.types import *
```

```
struct1 = StructType().add("placeid", IntegerType(), True).add("placename", StringType(), True)
```

```
foodplaces = spark.read.schema(struct1).csv('foodplaces213024.txt')
```

```
foodplaces.printSchema()
```

```
foodplaces.show(5)
```

```
>>> tab2=StructType().add("placeid",IntegerType(),True).add("placename", StringType(),True)
>>> foodplaces=spark.read.schema(tab2).csv('hdfs:///user/csp554/foodplaces213024.txt')
>>> foodplaces.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> foodplaces.show(5)
+-----+-----+
|placeid|  placename|
+-----+-----+
|      1| China Bistro|
|      2|   Atlantic|
|      3|  Food Town|
|      4|   Jake's|
|      5|  Soup Bowl|
+-----+-----+
```

Exercise 3

Step A

Register the DataFrames created in exercise 1 and 2 as tables called "foodratingsT" and "foodplacesT"

```
foodratings.createOrReplaceTempView("foodratingsT")
```

```
foodplaces.createOrReplaceTempView("foodplacesT")
```

```
>>> foodratings.createOrReplaceTempView("foodratingsT")
>>> foodplaces.createOrReplaceTempView("foodplacesT")
```

Step B

Use a SQL query on the table “foodratingsT” to create a new DataFrame called **foodratings_ex3** holding records which meet the following condition: **food2 < 25** and **food4 > 40**. Remember, when defining conditions in your code use maximum parentheses

```
foodratings_ex3 = spark.sql("SELECT * from foodratingsT where food2 < 25 and food4 > 40")
foodratings_ex3.printSchema()
```

```
foodratings_ex3.show(5)
```

```
>>> foodratings_ex3 = spark.sql("SELECT * from foodratingsT where food2 < 25 and food4 > 40")
```

```
|>>> foodratings_ex3.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)
```

```
>>> █
```

```
>>> foodratings_ex3.show(5)
```

```
+---+---+---+---+---+---+
|name|food1|food2|food3|food4|placeid|
+---+---+---+---+---+---+
[ | Mel| 42| 19| 35| 42| 3|
[ | Sam| 46| 1| 7| 45| 5|
[ | Sam| 50| 21| 48| 48| 3|
| Joy| 47| 2| 2| 49| 3|
| Joy| 31| 20| 33| 50| 1|
+---+---+---+---+---+---+
```

```
only showing top 5 rows
```

Step C

Use a SQL query on the table “foodplacesT” to create a new DataFrame called **foodplaces_ex3** holding records which meet the following condition: **placeid > 3**

```
foodplaces_ex3 = spark.sql("SELECT * from foodplacesT where placeid > 3")
foodplaces_ex3.printSchema()
```

```
foodplaces_ex3.show(5)
```

```
>>> foodplaces_ex3 = spark.sql("SELECT * from foodplacesT where placeid> 3")
>>> foodplaces_ex3.printSchema()
root
|-- placeid: integer (nullable = true)
|-- placename: string (nullable = true)
```

```
>>> foodplaces_ex3.show(5)
```

```
+-----+-----+
|placeid|placename|
+-----+-----+
|      4|  Jake's |
|      5|Soup Bowl|
+-----+-----+
```

Exercise 4

Use a transformation(not aSparkSQL query) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex4 that includes only those records (rows) where the 'name' field is "Mel" and food3 < 25

```
foodratings_ex4 = foodratings.filter(foodratings.name == "Mel").filter(foodratings.food3 < 25)
```

```
foodratings_ex4.printSchema()
```

```
foodratings_ex4.show(5)
```

```
>>> foodratings_ex4 = foodratings.filter(foodratings.name == "Mel").filter(foodratings.food3 < 25)
>>> foodratings_ex4.printSchema()
root
|-- name: string (nullable = true)
|-- food1: integer (nullable = true)
|-- food2: integer (nullable = true)
|-- food3: integer (nullable = true)
|-- food4: integer (nullable = true)
|-- placeid: integer (nullable = true)
```

```
>>> foodratings_ex4.show(5)
```

```
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
| Mel|    3|   15|   17|   39|      3|
| Mel|   17|    5|   21|   34|      4|
| Mel|   38|   33|   13|   18|      4|
| Mel|   34|   32|   15|    3|      5|
| Mel|   43|   26|   11|   38|      5|
+-----+-----+-----+-----+-----+
```

only showing top 5 rows

Exercise 5

Use a transformation (not aSparkSQL query) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex5 that includes only the columns (fields) 'name' and 'placeid'

```
foodratings_ex5 = foodratings.select(foodratings.name, foodratings.placeid)
```

```
foodratings_ex5.printSchema()
```

```
foodratings_ex5.show(5)
```

```
>>> foodratings_ex5 = foodratings.select(foodratings.name, foodratings.placeid)
>>> foodratings_ex5.printSchema()
root
 |-- name: string (nullable = true)
 |-- placeid: integer (nullable = true)
```

```
>>> foodratings_ex5.show(5)
+-----+-----+
|name|placeid|
+-----+-----+
| Sam|      2|
| Mel|      3|
| Sam|      2|
| Sam|      4|
| Joy|      1|
+-----+-----+
only showing top 5 rows
```

Exercise 6

Use a transformation (not aSparkSQL query) to create a new DataFrame called ex6 which is the inner join, on placeid, of the DataFrames 'foodratings' and 'foodplaces' created in exercises 1 and 2

```
ex6 = foodratings.join(foodplaces, foodratings.placeid ==
```

```
foodplaces.placeid,"inner").drop(foodratings.placeid)
```

```
ex6.printSchema()
```

```
ex6.show(5)
```

```
row(name= Joe , placeid=3/, row(name= Mel , placeid=4/, row(name= Joe , placeid=3/, row(name= Joe , placeid=2/, row(
>>> ex6 = foodratings.join(foodplaces, foodratings.placeid == foodplaces.placeid, "inner").drop(foodratings.placeid)
>>> ex6.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)
```

```
>>> ex6.show(5)
```

name	food1	food2	food3	food4	placeid	placename
Sam	33	26	38	8	2	Atlantic
Mel	3	15	17	39	3	Food Town
Sam	22	14	36	14	2	Atlantic
Sam	27	16	50	47	4	Jake's
Joy	8	24	1	11	1	China Bistro

```
only showing top 5 rows
```